



# Introduction to Artificial Intelligence (CS-487)

## Assignment #1

Due on Oct 29, 2022

*Professor I. Tsamardinos*

University of Crete  
Department of Computer Science

**Nikolaos Kougioulis** (ID 1285)

October 29, 2022

## Exercise 1

Both the performance measure and the utility function measure how well an agent does. Explain how they differ.

### Solution:

Indeed, both the performance measure and the utility function serve as a measure of how "well" an agent performs. There is however a slight difference, depending on which perspective we choose (the designer of the agent(a human) or the agent itself.

A performance measure is in the mind of the designer of the agent, or in the mind of the users of the agent. The performance of the agent is evaluated on how preferable their sequence of actions is from the human's point of view. As such, a performance measure is objective and can include information unavailable to the agent (like in a partially observable, sequential environment). The performance measure may be explicit or implicit (the agent may perform the correct task, but have no idea why). That's why we need to assume the performance measure can be specified correctly.

On the other hand, a utility function includes only information available to the agent itself, the states and sequence of actions. The utility function is an internalization of the performance measure itself and is ideally an estimate of the performance measure (if both the utility function and performance measure are in agreement, an agent choosing actions that maximize its utility function will be *rational* according the performance measure.)

## Exercise 2

Can such graph exist in which  $A^*$  extends more nodes than the Depth-First Search algorithm; If so, draw an example of such a graph. If not, explain why this is impossible.

### Solution:

We expect Depth-First Search to expand more nodes than  $A^*$  search with an admissible heuristic.

A lucky Depth-First Search algorithm may however expand fewer nodes than  $A^*$ , simply by expanding exactly those nodes on the optimal path to reach the goal state (if the optimal path consists of  $k$  nodes, then a lucky Depth-First Search may expand exactly those  $k$  nodes), without the need for backtracking (backing up to the next deepest node that has unexpanded successors).

So in the case of a lucky Depth-First Search,  $A^*$  could possibly expand more nodes before finding the optimal path. Moreover,  $A^*$  may expand more nodes than Depth-First Search with an inadmissible heuristic.

We illustrate an example similar to the routes of Romania problem with starting node A, final node D and heuristic function values  $h(A) = 500$ ,  $h(B) = 400$ ,  $h(C) = 100$ ,  $h(D) = 0$ ,  $h(E) = 330$ ,  $h(F) = 350$  (see 1).




Figure 1: Formulation of the state space, similar to the Romania problem, with initial state node A and final state node D.

In figure 2, Depth-First Search is illustrated. Starting with the initial node, Depth-First Search always expands the deepest node on the frontier. The algorithm is being lucky since it has reached the goal state expanding those nodes on the optimal path, without the need for backtracking (after reaching the deepest level of the tree, with no more successors, returning to the next deepest node on the frontier that has unexpanded successors.). We assume that in our implementation Depth-First Search takes care checking each nodes for cycles, as to not get stuck in infinite loops due to the cyclic state space.




Figure 2: First steps of Depth-First Search. Starting with the initial node A, the deepest node on the frontier is expanded until there are no more successors (deepest point), then returns to the next deepest node on the frontier with unexpanded successors. Depth-First Search is not cost-optimal, as it returns the first solution it finds, even if it is not the cheapest.

A-star starts by expanding the initial node and proceeds to evaluate each step, a Best-First Search algorithm using the evaluation function  $f(n) = g(n) + h(n)$  where  $g(n)$  is the path cost from the initial state to node




Figure 3: Final step of Depth-First Search, being lucky by reaching the goal state with the optimal path without backtracking and expanding more nodes.

$n$  (the costs for each available transition between states is marked between the edges in 1) and  $h(n)$  is the estimated cost of the shortest path from  $n$  to the goal state.

Notice that when the goal state  $D$  appears on the search tree (4), the A-star algorithm will not settle for a solution with such cost because the unexpanded node  $B$  at the frontier indicates that there might be a solution with lower cost through node  $B$ . That is when, in the next step, node  $B$  will be expanded and A-star will lead to the optimal solution.

As a result we see how in this example, due to the formulation of the problem and the heuristic function  $h$ , how Depth-First Search was lucky in obtaining the optimal path in fewer steps than A-star, although as we discussed in the first paragraph it is not usually the case. Also A-star is cost-optimal while Depth-First Search is not; it will return the first solution it finds even if not the cheapest.

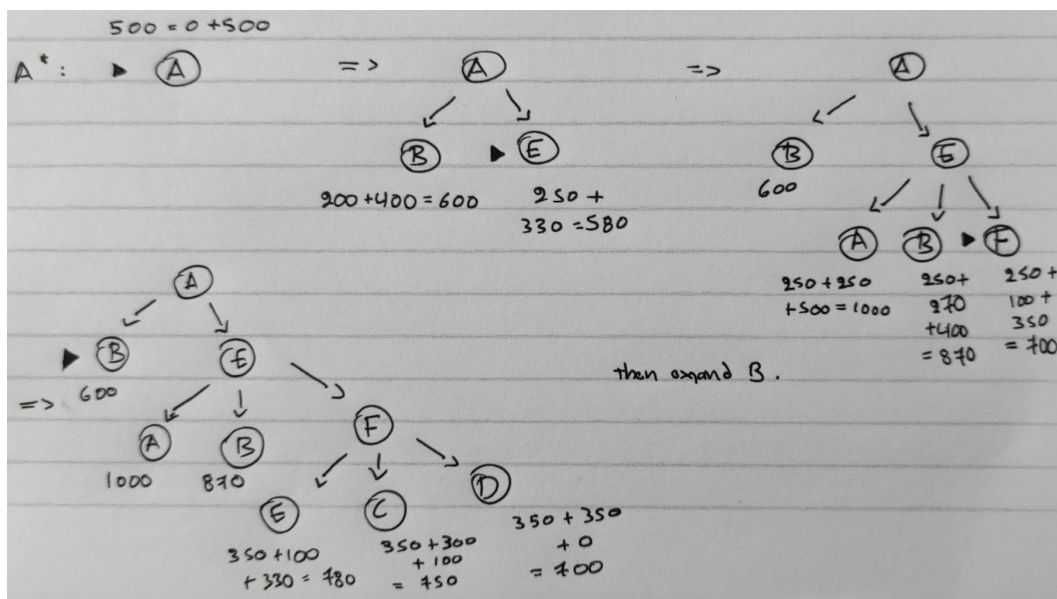


Figure 4: Progress of the  $A^*$  algorithm with the goal of reaching state  $D$ .

## Exercise 3

1. Suppose we run a greedy search algorithm with  $h(n) = -g(n)$ . What sort of search will the greedy search emulate?
2. Sometimes there is no good evaluation function for a problem, but there is a good comparison method: a way to tell if one node is better than another, without assigning numerical values to either. Show that this is enough to do a best-first search. What properties of best-first search do we give up if we only have a comparison method?

### Solution:

1. The best candidate nodes will be the ones with the longest cost paths, because the deeper the node the better its cost. So it emulates depth-first search (DFS).
2. Best-first search is implemented using a priority queue, so we can still do best-first search by sorting the queue only by comparison between the elements of the queue. But because there is no quantitative information about whether one node is better than another, we cannot combine the results of the comparison with a function like  $g(n)$  or  $h(n)$  so we cannot implement  $A^*$  search, so we give up completeness and optimality.

## Exercise 4

In this exercise you will run various informed and uninformed search algorithms to solve the 8-puzzle problem. You can use the code of the book given to you (Note that Assignment1.ipynb file need to be inside aima-search folder). The algorithms you can use can be found on search.py. You will have to run the code 100 times for the 8-puzzle, for random puzzles each time. Because some algorithms may take too long to solve a problem, we set an upper limit on the number of states that each algorithm can visit ( $10^7$  in our case). In case the algorithm fails to find a solution (it has reached the limit of the number of states it can visit, it has reached the limit of the memory it can use or for some other reason) the value  $-1$  is stored as the number of states and the size of the solution.

**Use at least 4 algorithms of your choice.**

You have to compare the algorithms based on the above measurements in terms of the complexity and quality of the solution, since a solution has been found. Also calculate how many times each algorithm found a solution. How you compare them is up to you. For example, you can calculate the transaction factor, various statistics (mean values, standard deviation, etc.), create graphs (see also in the book how search algorithms are compared). You can use whatever tools you want for this purpose and modify the given code accordingly. One should be able to decide which algorithm to use based on your comparison. Also answer the following questions:

- What results would you expect based on the theory;
- Are they verified by experiments? Comment on that.
- Which algorithm would you choose to use? Justify your answer.
- How does the size of the optimal solution affect the performance of the algorithms;

Experiment with the maximum number of states (variable max-actions) (i.e. run the algorithms several times with different variable values). What do you notice? Justify your answers.

### Solution:

For time complexity issues, do the nature of the time complexity of the algorithms, the number of max actions has been set to  $10^5$ , which still yields meaningful results on the performance of the search algorithms.

Used algorithms include A-star Search, Recursive Best-First Search (R-BFS), Depth-First Search (DFS) and Uniform-Cost Search.

```

1  names = ['A-star', 'Recursive BFS', 'Depth-First Search', 'Uniform-Cost Search']
2
3  count_values = [count_times, count1_times, count2_times, count3_times]
4  mean_values = [count_mean, count1_mean, count2_mean, count3_mean]
5  var_values = [count_var, count1_var, count2_var, count3_var]
6
7  plt.figure(figsize=(9, 4))
8
9  plt.scatter(x=range(1, count+1), y=count_times, label="A-star")
10 plt.scatter(x=range(1, count1+1), y=count1_times, label="Recursive BFS")
11 plt.scatter(x=range(1, count2+1), y=count2_times, label="Depth-First Search")
12 plt.scatter(x=range(1, count3+1), y=count3_times, label="Uniform-Cost Search")
13 plt.suptitle('Running Times of Successful Runs (in seconds)')
```

```

14     plt.xlabel("Run")
15     plt.ylabel("seconds")
16     plt.legend()
17     plt.show()
18

```

Listing 1: Python Code for scatterplot of the running times of Successful Runs (in seconds) for each Search Algorithm

We evaluate the performance of the presented algorithms on the basis of completeness, cost optimality, time and space complexity:

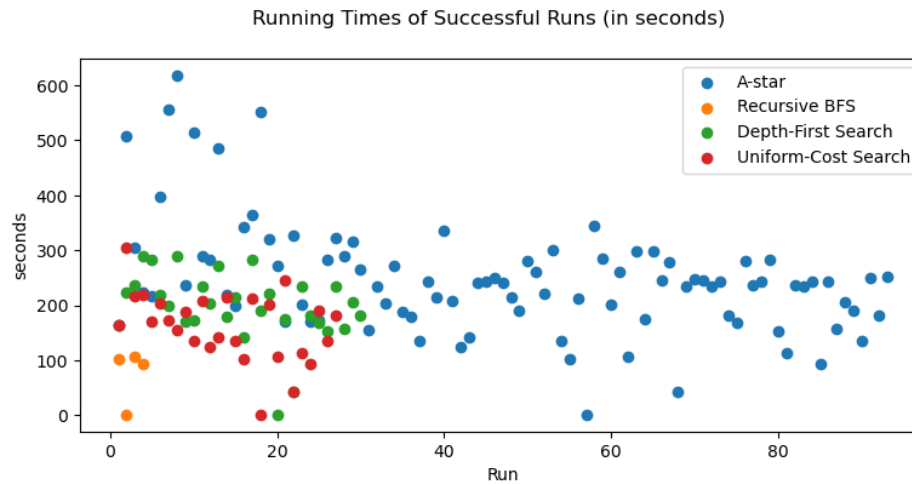


Figure 5: Scatterplot of Running Times for each solvable initial state

- A-star is complete and optimal, provided  $h(n)$  is an admissible heuristic (in our case, in the 8-puzzle problem, it is). The space complexity, which is exponential, is still an issue in practice), which we verify by our experiments.
- Recursive-BFS is a robust and optimal version of Best-First-Search that uses limited memory resources compared to standard BFS. So given enough time it would solve problems in which A-star runs out of memory.
- Depth-First Search expands the deepest node first, it is neither complete nor optimal but has linear space complexity.
- Uniform-Cost Search expands the nodes with the lowest  $g(n)$ , is optimal for general action costs but running time remains an issue.

```

1  #####
2  Total random initial states: 100
3  That are solvable: 100
4  A-star: 93 successes
5  Recursive BFS: 4 successes
6  Depth-First Search: 30 successes
7  Uniform-Cost Search: 27 successes
8  #####
9  A-star mean time: 246.57977447714856
10 Recursive BFS mean time: 75.70691114664078
11 Depth-First Mean time: 197.30689787069957

```

```

12 Uniform-Cost Search Mean time: 162.12999637921652
13 #####
14 A-star variance time: 10686.768832753627
15 Recursive BFS variance time: 2482.7021897812915
16 Depth-First Search variance time: 4076.8551931757356
17 Uniform-Cost Search variance time: 3998.614603300636
18 --- Elapsed Time: 26375.817984342575 seconds ---
19

```

Listing 2: Output of the 4 algorithms

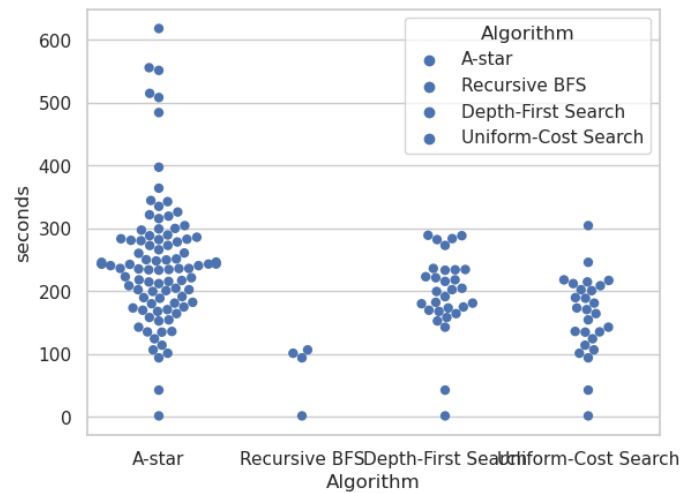


Figure 6: Swarmplot of the Running Times of Successful Runs (in seconds)




Figure 7: Running Time Sample Means of Successful Runs (in seconds)

We begin by testing 100 solvable initial random states, with an upper bound for the total number of actions. We then generate a scatterplot with the running times of the successful runs (the ones that reach a goal state) as well as a barplot with the mean and variance of each algorithm. A swarmplot (similar to a boxplot) is also generated, to give a broader picture of each sample distribution and outliers.

With the number of maximum actions we have specified, A-star has reached the goal state in almost every of the solvable initial states, in general using more running time but with much higher variance in the sample runs. RBFS has reached only a handful of them, while Depth-First and Uniform-Cost search have performed






Figure 8: Running Time Sample Variances of Successful Runs (in seconds)

in a very similar manner.

To assert whether the four algorithms perform different than each other, we conduct pairwise t-tests, on a significance level  $\alpha = 0.05$ , a total of  $\binom{4}{2} = 6$  times assuming False variances in all cases except with Depth-First Search and Uniform-Cost Search where our sample variances appear highly similar in value. The Hypothesis testing between the true means  $\mu_0$  and  $\mu_1$  of the two paired samples is:

$$H_0 : \mu_0 = \mu_1, H_1 : \mu_0 \neq \mu_1$$

```
1 print("Pairwise t-test for running times of Depth-First Search and Uniform-Cost Search:",
2 stats.ttest_ind(a=count2_times, b=count3_times, equal_var=True))
```

Listing 3: Example of Paired t-test between Depth-First Search and Uniform-Cost Search assuming equal variance in the distribution

Our significance level  $\alpha$  corresponds to Error type I which is the Probability we reject the null hypothesis if true. A p-value less than  $\alpha$  indicates strong evidence against the null hypothesis, hence the null hypothesis is not accepted. A p-value higher than  $\alpha$  indicates strong evidence supporting the null hypothesis, hence the null hypothesis is accepted.

For each test, we get a p-value of 0.0027, 0.0025, 1.8315, 0.009, 0.0297 and 0.04160 respectively. Hence A-star on average performs slower than DFS and Uniform-Cost Search, while for Recursive BFS our data does not give us a clear distinction between the running times of the two algorithms from our 100 solvable initial states.

Due to the informed nature of A-star, as well as the fact that it is complete and optimal, it would generally be our first choice for a Search Algorithm in practice. By expanding the allowed number of actions, A-star finds a solution to all 100 solvable initial states, while other algorithms waste enormous amounts of running time and failing to reach a goal state due to not being complete. By lowering the amount of maximum actions, most fail to reach the goal state, as expected, unless they get "lucky" with the initial state, requiring a few number of actions. Even in these cases, A-star finds most solutions rather quickly.

Full Code listing is available at the accompanied .ipynb Jupyter Notebook, as well as at the appendix of this document.

```

1     from search import * #import all search algorithms
2     import time
3     import random
4     import timeit
5
6     import matplotlib.pyplot as plt #use matplotlib for easthetic plots
7     import seaborn as sns #for swarmplot
8     import numpy as np
9     import pandas as pd
10
11     import statistics #for mean and variance
12     import scipy.stats as stats #for hypothesis testing
13
14     class EightPuzzle(Problem):
15         """ The problem of sliding tiles numbered from 1 to 8 on a 3x3 board, where one of
16         the
17         squares is a blank. A state is represented as a tuple of length 9, where element at
18         index i represents the tile number at index i (0 if it's an empty square) """
19
20         max_actions = 10**5 #limit for number of actions
21         current_actions = 0
22
23         def __init__(self, initial, goal=(1, 2, 3, 4, 5, 6, 7, 8, 0)):
24             """ Define goal state and initialize a problem """
25             super().__init__(initial, goal)
26
27         def find_blank_square(self, state):
28             """Return the index of the blank square in a given state"""
29
30             return state.index(0)
31
32         def actions(self, state):
33             """ Return the actions that can be executed in the given state.
34             The result would be a list, since there are only four possible actions
35             in any given state of the environment """
36
37             possible_actions = ['UP', 'DOWN', 'LEFT', 'RIGHT']
38             index_blank_square = self.find_blank_square(state)
39
40             if index_blank_square % 3 == 0:
41                 possible_actions.remove('LEFT')
42             if index_blank_square < 3:
43                 possible_actions.remove('UP')
44             if index_blank_square % 3 == 2:
45                 possible_actions.remove('RIGHT')
46             if index_blank_square > 5:
47                 possible_actions.remove('DOWN')
48
49             return possible_actions
50
51         def result(self, state, action):
52             """ Given state and action, return a new state that is the result of the action.
53             Action is assumed to be a valid action in the state """
54
55             # blank is the index of the blank square
56             blank = self.find_blank_square(state)
57             new_state = list(state)
58
59             delta = {'UP': -3, 'DOWN': 3, 'LEFT': -1, 'RIGHT': 1}
60             neighbor = blank + delta[action]
61             new_state[blank], new_state[neighbor] = new_state[neighbor], new_state[blank]
62             self.current_actions +=1

```

```

62         return tuple(new_state)
63
64     def goal_test(self, state):
65         """ Given a state, return True if state is a goal state or False, otherwise """
66
67         return (state == self.goal) or (self.current_actions > self.max_actions)
68
69     def check_solvability(self, state):
70         """ Checks if the given state is solvable """
71
72         inversion = 0
73         for i in range(len(state)):
74             for j in range(i + 1, len(state)):
75                 if (state[i] > state[j]) and state[i] != 0 and state[j] != 0:
76                     inversion += 1
77
78         return inversion % 2 == 0
79
80     def h(self, node):
81         """ Return the heuristic value for a given state. Default heuristic function
82         used is
83         h(n) = number of misplaced tiles """
84
85         return sum(s != g for (s, g) in zip(node.state, self.goal))
86
87     def run_singleTest(initial_state):
88
89         puzzle = EightPuzzle(initial_state)
90         results = []
91
92         puzzle.current_actions = 0
93
94         # start timing your execution
95         t0 = time.time()
96
97         # save important parameters
98         algorithm_results = {'Algorithm': 'A-star', #informed search
99                             'Initial State': initial_state,
100                             'Final_State': astar_search(puzzle)}
101
102         t1 = time.time() - t0
103         algorithm_results['Time'] = t1
104         results.append(algorithm_results)
105
106         #####
107         puzzle.current_actions = 0
108         # start timing your execution
109         t0 = time.time()
110         # save important parameters
111         algorithm_results = {'Algorithm': 'RBFS', #uninformed
112                             'Initial State': initial_state, #u
113                             'Final_State': recursive_best_first_search(puzzle)}
114
115         t1 = time.time() - t0
116         algorithm_results['Time'] = t1
117         results.append(algorithm_results)
118
119         #####
120         puzzle.current_actions = 0
121         # start timing your execution
122         t0 = time.time()
123         # save important parameters

```

```

123     algorithm_results = {'Algorithm': 'Depth-First Search', #uninformed
124                          'Initial State': initial_state,
125                          'Final_State': depth_first_graph_search(puzzle)} #breadth_first_search
    very slow to converge
126
127     t1 = time.time() - t0
128     algorithm_results['Time'] = t1
129     results.append(algorithm_results)
130
131     #####
132     puzzle.current_actions = 0
133     # start timing your execution
134     t0 = time.time()
135     # save important parameters
136     algorithm_results = {'Algorithm': 'Uniform-Cost Search', #very slow to converge
137                          'Initial State': initial_state,
138                          'Final_State': uniform_cost_search(puzzle)}
139     t1 = time.time() - t0
140     algorithm_results['Time'] = t1
141     results.append(algorithm_results)
142
143     return results
144
145     t0 = time.time()
146
147     goal = (1, 2, 3, 4, 5, 6, 7, 8, 0)
148     print('Goal:', goal)
149
150     size = 10**2 #total number of initial states
151
152     count = count1 = count2 = count3 = 0
153     total_solvable = 0
154
155     #running times for each initial solution (in seconds)
156     count_times = []
157     count1_times = []
158     count2_times = []
159     count3_times = []
160
161     i=0
162
163     #while (i < size+1): #test any initial state
164     while(total_solvable < 100): #test solvable initial states only
165         initial_state = tuple(random.sample(list(goal),9))
166
167         #check if chosen initial state is solvable
168         solvable = EightPuzzle(initial_state).check_solvability(initial_state)
169
170         if(not solvable):
171             print('Initial State unsolvable, moving to the next one:')
172         else:
173             total_solvable +=1;
174             t_solve0 = time.time()
175
176             results = run_singleTest(initial_state)
177             print('Initial State', i , ':', str(initial_state))
178
179             for result in results:
180                 if result['Final_State'].state == goal:
181                     if(str(result['Algorithm']) == 'A-star'):
182                         count+=1
183                         count_times.append(time.time() - t_solve0)

```

```

184         elif(str(result['Algorithm']) == 'RBFS'):
185             count1+=1
186             count1_times.append(time.time() - t_solve0)
187         elif(str(result['Algorithm']) == 'Depth-First Search'):
188             count2+=1
189             count2_times.append(time.time() - t_solve0)
190         elif(str(result['Algorithm']) == 'Uniform-Cost Search'):
191             count3+=1
192             count3_times.append(time.time() - t_solve0)
193
194         print('Algorithm: ' + str(result['Algorithm']) + ' succeeded , Execution
Time: ' + str(round(result['Time'],2)), 's')
195         print('Resulting solution: ' + str(result['Final_State'].solution()) + '
ending in board: ' + str(result['Final_State'].state))
196         print('-----')
197         else:
198             print('Algorithm: ' + str(result['Algorithm']) + ' failed , Execution Time
: ' + str(round(result['Time'],2)), 's')
199             print('Resulting solution: ' + str(result['Final_State'].solution()) + '
ending in board: ' + str(result['Final_State'].state))
200             print('-----')
201     i=i+1
202
203     count_mean = count1_mean = count2_mean = count3_mean = 0
204
205     count_var = count1_var = count2_var = count3_var = 0
206
207     if(len(count_times) > 0):
208         count_mean = statistics.mean(count_times)
209         count_var = statistics.variance(count_times)
210
211     if(len(count1_times) > 0):
212         count1_mean = statistics.mean(count1_times)
213         count1_var = statistics.variance(count1_times)
214
215     if(len(count2_times) > 0):
216         count2_mean = statistics.mean(count2_times)
217         count2_var = statistics.variance(count2_times)
218
219     if(len(count3_times) > 0):
220         count3_mean = statistics.mean(count3_times)
221         count3_var = statistics.variance(count3_times)
222
223     print('##=====##')
224     #print('Max Actions:', max_actions)
225     print('Total random initial states:', size)
226     print('That are solvable:', total_solvable)
227     print('A-star:', count, 'successes')
228     print('Recursive BFS:', count1, 'successes')
229     print('Depth-First Search:', count2, 'successes')
230     print('Uniform-Cost Search: ', count3, 'successes')
231     print('##=====##')
232     print('A-star mean time:', count_mean)
233     print('Recursive BFS mean time:', count1_mean)
234     print('Depth-First Mean time:', count2_mean)
235     print('Uniform-Cost Search Mean time:', count3_mean)
236     print('##=====##')
237     print('A-star variance time:', count_var)
238     print('Recursive BFS variance time:', count1_var)
239     print('Depth-First variance time:', count2_var)
240     print('Uniform-Cost variance time:', count3_var)
241     print("--- Elapsed Time: %s seconds ---" % (time.time() - t0))

```

```

242
243 names = ['A-star', 'Recursive BFS', 'Depth-First Search', 'Uniform-Cost Search']
244
245 count_values = [count_times, count1_times, count2_times, count3_times]
246 mean_values = [count_mean, count1_mean, count2_mean, count3_mean]
247 var_values = [count_var, count1_var, count2_var, count3_var]
248
249 plt.figure(figsize=(9, 4))
250
251 plt.scatter(x=range(1, count+1), y=count_times, label="A-star")
252 plt.scatter(x=range(1, count1+1), y=count1_times, label="Recursive BFS")
253 plt.scatter(x=range(1, count2+1), y=count2_times, label="Depth-First Search")
254 plt.scatter(x=range(1, count3+1), y=count3_times, label="Uniform-Cost Search")
255 plt.suptitle('Running Times of Successful Runs (in seconds)')
256 plt.xlabel("Run")
257 plt.ylabel("seconds")
258 plt.legend()
259 plt.show()
260
261 sns.set(style="whitegrid")
262 sns.swarmplot(x="Algorithm", y="seconds", data=pd.DataFrame({"seconds": count_values[0],
263     "Algorithm": names[0]}), size=6, hue='Algorithm')
264 sns.swarmplot(x="Algorithm", y="seconds", data=pd.DataFrame({"seconds": count_values[1],
265     "Algorithm": names[1]}), size=6, hue='Algorithm')
266 sns.swarmplot(x="Algorithm", y="seconds", data=pd.DataFrame({"seconds": count_values[2],
267     "Algorithm": names[2]}), size=6, hue='Algorithm')
268 sns.swarmplot(x="Algorithm", y="seconds", data=pd.DataFrame({"seconds": count_values[3],
269     "Algorithm": names[3]}), size=6, hue='Algorithm')
270 plt.show()
271
272 plt.figure(figsize=(9, 4))
273
274 plt.bar(names, mean_values, label="Average Mean ")
275 plt.suptitle('Running Time Sample Means of Successful Runs (in seconds)')
276 plt.legend()
277 plt.show()
278
279 plt.bar(names, var_values, label= names)
280 plt.suptitle('Running Time Sample Variances of Successful Runs (in seconds)')
281 plt.legend()
282
283 print("Pairwise t-test for running times of A-star and RBFS",
284 stats.ttest_ind(a=count_times, b=count1_times, equal_var=False))
285
286 print("Pairwise t-test for running times of A-star and Depth-First Search:",
287 stats.ttest_ind(a=count_times, b=count2_times, equal_var=False))
288
289 print("Pairwise t-test for running times of A-star and Uniform-Cost Search:",
290 stats.ttest_ind(a=count_times, b=count3_times, equal_var=False))
291
292 print("Pairwise t-test for running times of RBFS and Depth-First Search:",
293 stats.ttest_ind(a=count1_times, b=count2_times, equal_var=False))
294
295 print("Pairwise t-test for running times of RBFS and Uniform-Cost Search:",
296 stats.ttest_ind(a=count1_times, b=count3_times, equal_var=False))
297
298 print("Pairwise t-test for running times of Depth-First Search and Uniform-Cost Search:",
299 stats.ttest_ind(a=count2_times, b=count3_times, equal_var=True))
300

```

Listing 4: Full Code Listing