



Introduction to Artificial Intelligence (CS-487)

Pac-man Projects Contest (Phase C)

Due on Jan 24, 2023

Professor I. Tsamardinos

University of Crete
Department of Computer Science

Nikolaos-Modestos Kougioulis (ID 1285)

January 20, 2023

Introduction

For Phase C of the Pac-Man Projects we are asked to create an agent that successfully competes and solves the original map of Pac-Man (available with the argument `-l originalClassic`), in the form of competition with fellow students. The implementation should not be laggy and require an extensive amount of time to complete (such as looping back and forth to finite states), as well as being advised to make use of the Expectimax agent that was created for the needs of Phase B of the Projects.

We combine properties of the Expectimax, FoodSearch and Reflex agents to create a new agent called the *Hybrid Agent*, that consistently wins high scoring games in a time-efficient manner.

Contents

Introduction	2
Methods - The <i>HybridAgent</i> class	3
Results	5
Discussion	6
Appendix	7

List of Figures

1	Gameplay of the Pac-Man Projects on <i>originalClassic</i> map. The screenshot illustrates our <i>HybridAgent</i> mid-game.	4
2	Scores on 10 iterations at the <i>originalClassic</i> map of our <i>HybridAgent</i> , by running <code>python3 pacman.py -p HybridAgent -a depth=3 -l originalClassic -n 10</code>	5
3	Scores on 10 iterations at the <i>originalClassic</i> map of an average human player, running <code>python3 pacman.py -l originalClassic -n 10</code>	5

Methods - The *HybridAgent* class

In this section we introduce a new agent class to solve the task of the project. Before coming up with a non-naive and trivial approach to the problem of solving the *originalClassic* map, we must make the following assumptions:

- **Assumption 1:** If implementing Expectimax, always assume depth 3. This derives from the test command provided in the guidelines¹.
- **Assumption 2:** We are not restricted to a single agent behavior but may apply collection of agent behaviors to a single agent, hereon called the *Hybrid Agent*.
- **Assumption 3:** Assume we can evaluate the performance of the hybrid agent by comparing against an average human player (and surely not against Billy Mitchell).

For our implementation, we follow an approach similar to what an average human player would follow. By average we refer to an individual familiar with the rules and goals of Pac-Man, but not knowledgeable with how each individual ghost behaves and which actions prove optimal in the long run, maximizing the final score. We use the following three rules:

Rule 1: If the closest distance to a ghost is over k , then find the optimal path from the current position to the closest food pellet. That is, if the following is satisfied

$$\min(\{d_M(x_p, x_i), i = 1, 2, 3\}) > k \Rightarrow \min(\{|x_p^1 - x_i^1| + |y_p^1 - y_i^2|, i = 1, 2, 3\}) > k$$

proceed with a food pellet searching algorithm, where x_i are the three adversarial of Pac-Man: Blinky (red), Inky (cyan), Pinky (pink) and Clyde (orange).

Rule 2: If the closest distance to a ghost is at most k , then perform Expectimax assuming a uniform distribution for the adversarial's actions.

Rule 3: In case a food pellet is unreachable, handle illegal actions by randomly being a left turn only or right turn only agent.

We define *HybridAgent* as the agent combining the above three rules. The intuition behind this implementation is simple: If ghosts are far from Pac-Man, then they do not pose a realistic threat, so it is sensible to ignore them altogether and focus on collecting the most amount of food pellets. However when ghosts get close to Pac-Man to a constant k , then Pac-Man can sense death is near and uses Expectimax to avoid the ghosts while collecting as much food pellets as possible. The evaluation function has also been altered a bit in order to eat a scared ghost when eating a power capsule (by returning $+\infty$ as score) and run away from certain death (by returning $-\infty$ as score).

The path to the nearest food pellet is obtained using Breadth-First Search: At each step of the game we obtain the coordinates of Pac-Man and then proceed to get the coordinates of each legal action. This is better illustrated with the following code snippet:

```

1  #use Breadth-First Search to find the path to the closest food
2
3  startingNode = curr_pos
4  queue = util.Queue()
```

¹we are recommended to test our implementation by running `pacman.py -p ExpectimaxAgent -a depth=3 -l originalClassic`




Figure 1: Gameplay of the Pac-Man Projects on *originalClassic* map. The screenshot illustrates our *HybridAgent* mid-game.

```

5  visited = set() #use dictionary for faster lookup
6
7  queue.push((startingNode, []))
8
9  if curr_pos == closest_food:
10     return []
11
12  #return the first action that leads to the closest food
13  while not queue.isEmpty():
14     node, path = queue.pop()
15     if node not in visited:
16         visited.add(node)
17     if node == closest_food:
18         print("Path to follow:", path)
19         return path[0]
20  for action in legal:
21     if action == Directions.STOP:
22         continue
23     #get coordinates of the next nodes in the path
24     if action == Directions.NORTH:
25         successor = (node[0], node[1] + 1)
26     elif action == Directions.SOUTH:
27         successor = (node[0], node[1] - 1)
28     elif action == Directions.WEST:
29         successor = (node[0] - 1, node[1])
30     elif action == Directions.EAST:
31         successor = (node[0] + 1, node[1])
32     elif action == Directions.STOP:
33         successor = node
34     #if none of the above, then the action is not valid
35     else:
36         continue
37     if successor not in visited and successor not in walls:
38         queue.push((successor, path + [action]))
39

```

Full code of the DFS implementation and the HybridAgent class is available in the appendix of this document, as well on the associated python source file.

Results

As per guidelines, *HybridAgent* is run on the *originalClassic* layout 10 times. The scores of each game, the average score and win rate is shown in Figure 2. We implement our HybridAgent using $k = 4$. The scores of 10 human played game instances are illustrated in Figure 3.

Our HybridAgent wins all 10 games, with a 3112.9 average score, while for the human player 2 are lost.

To compare the human player with HybridAgent, we perform the Wilcoxon Signed-rank test, a non-parametric statistical test, for sample size $n = 10$. For each instance, we compute the difference between the two scores D_i . We then obtain the absolute values $|D_1|, |D_2|, \dots, |D_n|$ and order them from smallest to largest by assigning them ranks from 1 to n , $r(|D_1|), r(|D_2|), \dots, r(|D_n|)$. We also keep a record of the original signs of the differences, notating I^+ and I^- the list of indices i for which the signs were positive and negative respectively. The Wilcoxon statistic W is the smallest of W^+ (the sum of the positive ranks) and W^- (the sum of the negative ranks),

$$W = \min(W^+, W^-) = \min\left(\sum_{i \in I^+} r(|D_i|), \sum_{i \in I^-} r(|D_i|)\right)$$

which does not depend on the distribution of the samples (hence non-parametric). For large n , the W -statistic converges to the normal distribution. The null hypothesis is $H_0 : P(X_\alpha) > P(X_\beta) = 1/2$, that is, a randomly chosen observation from the first group is equally probable of being larger than a randomly drawn observation from the second group. Performing a two-tailed test, we obtain a mean difference -283.1 , sum of positive ranks $W^+ = 37$ and sum of negative ranks $W^- = 18$. The Wilcoxon test-statistic is therefore $W = 18$. The critical value of W , obtained from a table of critical values for $n = 10$ and $\alpha = 0.05$ is 8. Since $W > 8$ we do not have statistically significant evidence to reject then null hypothesis. Hence the results are not statistically significant, so we can assume our HybridAgent performs as well as the average human player.

```
Pacman emerges victorious! Score: 2921
Average Score: 3112.9
Scores: 3427.0, 3502.0, 3138.0, 3324.0, 3302.0, 2723.0, 3214.0, 2711.0, 2867.0, 2921.0
Win Rate: 10/10 (1.00)
Record: Win, Win, Win, Win, Win, Win, Win, Win, Win, Win
nikolas@ideapad-3:~/Downloads/cs-487/Project Phase C/multiagent$
```

Figure 2: Scores on 10 iterations at the *originalClassic* map of our *HybridAgent*, by running `python3 pacman.py -p HybridAgent -a depth=3 -l originalClassic -n 10`.

```
Pacman emerges victorious! Score: 2833
Average Score: 2528.2
Scores: 3101.0, 3396.0, 3032.0, 152.0, 3245.0, 2863.0, 3377.0, 2978.0, 305.0, 2833.0
Win Rate: 8/10 (0.80)
Record: Win, Win, Win, Loss, Win, Win, Win, Win, Loss, Win
nikolas@ideapad-3:~/Downloads/cs-487/Project Phase C/multiagent$
```

Figure 3: Scores on 10 iterations at the *originalClassic* map of an average human player, running `python3 pacman.py -l originalClassic -n 10`.

Discussion

One computational improvement on the above implementation is to replace Breadth-First Search with an informed search algorithm, such as uniform-cost search or A-star search.

Ideally in larger maps, one can implement Markov Decision Processes with rewards. At each time step, the Markov process is in some state s and associates each action a resulting in a new state s' with a reward $R_a(s, s')$. These processes, as the name suggests, satisfy the Markov property known from Markov chains: the next state s' depends only on the current state s and the action a .

Appendix

```

1 class HybridAgent(MultiAgentSearchAgent):
2     """
3     Hybrid Agent
4     """
5
6     def getAction(self, gameState):
7
8         #for pacman
9         def max_value(gameState, depth):
10             legalActions = gameState.getLegalActions(0)
11             if not legalActions or depth == self.depth:
12                 return self.evaluationFunction(gameState)
13             v = float("-inf")
14             v = max(exp_value(gameState.generateSuccessor(0, action), 0 + 1, depth + 1) for
15 action in legalActions)
16             return v
17
18         #for all ghosts
19         def exp_value(gameState, agentIndex, depth):
20             legalActions = gameState.getLegalActions(agentIndex)
21             if not legalActions:
22                 return self.evaluationFunction(gameState)
23
24             prob = 1.0 / len(legalActions)
25             v = 0
26             for action in legalActions:
27                 newState = gameState.generateSuccessor(agentIndex, action)
28                 if agentIndex == gameState.getNumAgents() - 1:
29                     v += max_value(newState, depth) * prob
30                 else:
31                     v += exp_value(newState, agentIndex + 1, depth) * prob
32             return v
33
34         ghostPositions = gameState.getGhostPositions()
35         distanceToGhost = [util.manhattanDistance(gameState.getPacmanPosition(), ghost) for
36 ghost in ghostPositions]
37
38         if(min(distanceToGhost) > 4):
39
40             curr_pos = gameState.getPacmanPosition()
41             foods = gameState.getFood().asList() #list of food coordinates
42             capsules = gameState.getCapsules() #list of capsule coordinates
43             walls = gameState.getWalls().asList()
44             #closest food coordinate
45             closest_food = min(foods, key=lambda x: util.manhattanDistance(curr_pos, x))
46
47             legal = gameState.getLegalActions(0)
48
49             #return the action to the closest food dot
50             print("Closest food pellet: ", str(closest_food))
51             print("Legal actions: ", legal)
52
53             #use Breadth-First Search to find the path to the closest food
54             startingNode = curr_pos
55             queue = util.Queue()
56             visited = set() #use dictionary for faster lookup
57
58             queue.push((startingNode, []))

```

```

58     print("Current position: ", str(curr_pos))
59
60     if curr_pos == closest_food:
61         return []
62
63     #return the first action that leads to the closest food
64     while not queue.isEmpty():
65         node, path = queue.pop()
66         if node not in visited:
67             visited.add(node)
68             if node == closest_food:
69                 print("Path to follow:", path)
70                 return path[0]
71             for action in legal:
72                 if action == Directions.STOP:
73                     continue
74
75                 #get coordinates of the next nodes in the path
76                 if action == Directions.NORTH:
77                     successor = (node[0], node[1] + 1)
78                 elif action == Directions.SOUTH:
79                     successor = (node[0], node[1] - 1)
80                 elif action == Directions.WEST:
81                     successor = (node[0] - 1, node[1])
82                 elif action == Directions.EAST:
83                     successor = (node[0] + 1, node[1])
84                 elif action == Directions.STOP:
85                     successor = node
86                 #if none of the above, then the action is not valid
87                 else:
88                     continue
89                 if successor not in visited and successor not in walls:
90                     queue.push((successor, path + [action]))
91
92     #handling illegal moves if the closest food is not reachable by becoming a right
    turn or left turn reflex agent
93     if not gameState.hasFood(gameState.getPacmanPosition()[0] + 1, gameState.
    getPacmanPosition()[1]) and gameState.hasFood(gameState.getPacmanPosition()[0] - 1,
    gameState.getPacmanPosition()[1]):
94         legal = gameState.getLegalActions(0)
95         current = gameState.getPacmanState().configuration.direction
96         if current == Directions.STOP: current = Directions.NORTH
97         left = Directions.LEFT[current]
98         if left in legal: return left
99         if current in legal: return current
100        if Directions.RIGHT[current] in legal: return Directions.RIGHT[current]
101        if Directions.LEFT[left] in legal: return Directions.LEFT[left]
102        return Directions.STOP
103    elif not gameState.hasFood(gameState.getPacmanPosition()[0] - 1, gameState.
    getPacmanPosition()[1]) and gameState.hasFood(gameState.getPacmanPosition()[0] + 1,
    gameState.getPacmanPosition()[1]):
104        legal = gameState.getLegalActions(0)
105        current = gameState.getPacmanState().configuration.direction
106        if current == Directions.STOP: current = Directions.NORTH
107        right = Directions.RIGHT[current]
108        if right in legal: return right
109        if current in legal: return current
110        if Directions.LEFT[current] in legal: return Directions.LEFT[current]
111        if Directions.RIGHT[right] in legal: return Directions.RIGHT[right]
112        return Directions.STOP
113    else:
114        if random.uniform(0, 1) < 0.5:

```



```
115         legal = gameState.getLegalActions(0)
116         current = gameState.getPacmanState().configuration.direction
117         if current == Directions.STOP: current = Directions.NORTH
118         right = Directions.RIGHT[current]
119         if right in legal: return right
120         if current in legal: return current
121         if Directions.LEFT[current] in legal: return Directions.LEFT[current]
122         if Directions.RIGHT[right] in legal: return Directions.RIGHT[right]
123         return Directions.STOP
124     else:
125         legal = gameState.getLegalActions(0)
126         current = gameState.getPacmanState().configuration.direction
127         if current == Directions.STOP: current = Directions.NORTH
128         left = Directions.LEFT[current]
129         if left in legal: return left
130         if current in legal: return current
131         if Directions.RIGHT[current] in legal: return Directions.RIGHT[current]
132         if Directions.LEFT[left] in legal: return Directions.LEFT[left]
133         return Directions.STOP
134
135     #if a ghost is close, death is near so use Expectimax
136     else:
137         best_action = None
138         legalActions = gameState.getLegalActions()
139         best_action = max(legalActions, key=lambda action: exp_value(gameState.
generateSuccessor(0, action), 1, 1))
140         print("A ghost is close! Performing Expectimax at location:", gameState.
getPacmanPosition())
141         return best_action
```

Listing 1: The HybridAgent class